

```
1 import pickle
2 import time
3
4 from sklearn.feature_extraction.text import TfidfVectorizer
5 from sklearn.decomposition import TruncatedSVD
6 from sklearn.pipeline import make_pipeline
7 from sklearn.preprocessing import Normalizer
8 from sklearn.neighbors import NearestNeighbors
9 from sklearn.neighbors.classification import KNeighborsClassifier
```



## NeighborsClassifier

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount()

```
1 X_train = []
2 y_train = []
3 X_test = []
4 y_test = []
5 for files in os.listdir(datasetPath):
6     for f in os.listdir(datasetPath+files):
7         d = open(datasetPath+files+"/"+f, "r")
8         data = str(d.read())
9         # Split the data into training and testing sets
```

```

9     label = re.searchn("[\d]+_(\d+)", r).group(1)
10    if files == "train":
11        #print(f)
12        X_train.append(data)
13        if (int(label) == 1):
14            y_train.append(1)
15        if (int(label) == 2):
16            y_train.append(1)
17        if (int(label) == 3):
18            y_train.append(2)
19        if (int(label) == 4):
20            y_train.append(2)
21        if (int(label) == 7):
22            y_train.append(3)
23        if (int(label) == 8):
24            y_train.append(3)
25        if (int(label) == 9):
26            y_train.append(4)
27        if (int(label) == 10):
28            y_train.append(4)
29
30
31    if files == "test":
32        #print(f)
33        X_test.append(data)
34        if (int(label) == 1):
35            y_test.append(1)

```

Saved successfully!



```

39        y_test.append(2)
40        if (int(label) == 4):
41            y_test.append(2)
42        if (int(label) == 7):
43            y_test.append(3)
44        if (int(label) == 8):
45            y_test.append(3)
46        if (int(label) == 9):
47            y_test.append(4)
48        if (int(label) == 10):
49            y_test.append(4)

```

```

1 X_train_raw = np.array(X_train)
2 y_train = np.array(y_train)
3 X_test_raw = np.array(X_test)
4 y_test = np.array(y_test)
5 target_names = []
6 labels = np.unique(y_train)
7 for label in labels:
8     target_names.append('Rating '+str(label))
9 print(X_train_raw.shape)

```

```
10 print(y_train.shape)
11 print(X_test_raw.shape)
12 print(y_test.shape)
```

```
(25000,)
(25000,)
(25000,)
(25000,)
```

```
1 import nltk
2 nltk.download('wordnet')
3 from nltk.stem import WordNetLemmatizer
4
5 def clean_dataset(X):
6     documents = []
7     stemmer = WordNetLemmatizer()
8     for sen in range(0, len(X)):
9
10         # Remove all the special characters
11         document = re.sub(r'\W', ' ', str(X[sen]))
12
13         # remove all single characters
14         document = re.sub(r'\s+[a-zA-Z]\s+', ' ', document)
15
16         # Remove single characters from the start
17         document = re.sub(r'\^[a-zA-Z]\s+', ' ', document)
18
19         # Removing trailing spaces with single space
20         document = re.sub(r'\s+', ' ', document, flags=re.I)
21
22         # Removing prefixed 'b'
23         document = re.sub(r'^b\s+', '', document)
24
25         # Converting to Lowercase
26         document = document.lower()
27
28         # Lemmatization
29         document = document.split()
30
31         document = [stemmer.lemmatize(word) for word in document]
32         document = ' '.join(document)
33
34         documents.append(document)
35     return documents
36 X_train_raw = clean_dataset(X_train_raw)
37 X_test_raw = clean_dataset(X_test_raw)
```

```
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
```

```
1 # Tfidf vectorizer:
```

```

2 vectorizer = TfidfVectorizer(max_df=0.7, max_features=500,
3                             min_df=5, stop_words='english',
4                             use_idf=True)
5
6 # Build the tfidf vectorizer from the training data ("fit"), and apply it
7 # ("transform").
8 X_train_tfidf = vectorizer.fit_transform(X_train_raw)
9
10 print(" Actual number of tfidf features: %d" % X_train_tfidf.get_shape()[1])
11
12 print("\nPerforming dimensionality reduction using LSA")
13 t0 = time.time()
14
15 # Project the tfidf vectors onto the first N principal components.
16 # Though this is significantly fewer features than the original tfidf vector,
17 # they are stronger features, and the accuracy is higher.
18 svd = TruncatedSVD(100)
19 lsa = make_pipeline(svd, Normalizer(copy=False))
20
21 # Run SVD on the training data, then project the training data.
22 X_train_lsa = lsa.fit_transform(X_train_tfidf)
23
24 print(" done in %.3fsec" % (time.time() - t0))
25
26 # explained_variance = svd.explained_variance_ratio_.sum()
27 # print(" Explained variance of the SVD step: {}".format(int(explained_variance * 100)))
28

```

Saved successfully!



o the test data as well.

```

31 X_test_tfidf = vectorizer.transform(X_test_raw)
32 X_test_lsa = lsa.transform(X_test_tfidf)

```

Actual number of tfidf features: 500

Performing dimensionality reduction using LSA  
done in 2.065sec

```

1 #scaling vector between 0 - 1
2 scaler = MinMaxScaler()
3 scaler.fit(X_train_lsa)
4 X_train_lsa_scaled =scaler.transform(X_train_lsa)
5
6 scaler = MinMaxScaler()
7 scaler.fit(X_test_lsa)
8 X_test_lsa_scaled =scaler.transform(X_test_lsa)

```

```

1 n_rangeK = 0
2 tMean = 0
3 n(array,column,folds=5):
4 = 0
5 i in range(1,folds+1):

```

```

6 n = sum + float(array[i][column])
7 n round(sum / folds,3)
8
9 formatMean(array):
10 al highestMean,bestAverageK
11 }
12 attedMean = ["Mean"]
13 i in range(1,20):
14 = mean(resultTable,i)
15
16 mattedMean.append(m)
17 highestMean < m:
18 highestMean = m
19 bestAverageK = i
20 n formattedMean
21 Table = [[]]
22
23 Fold(n_splits=5, random_state=None, shuffle=False)
24 in range(kf.get_n_splits(X_train_tfidf)):
25 train_index, test_index in kf.split(X_train_tfidf):
26 split_input_train_dataset , split_label_train_dataset = X_train_tfidf[train_index],y_train[
27 split_input_test_dataset , split_label_test_dataset = X_train_tfidf[test_index],y_train[tes
28
29 ltRow = [j+1]
30 i in range(1 , 20):
31 knn = KNeighborsClassifier(n_neighbors=i)
32 knn.fit(split_input_train_dataset , split_label_train_dataset)
33 cy_score(split_label_test_dataset,knn.predict(split_input_
34
35
36 late average for all value of K
37
38 Table.append(formatMean(resultTable))
39 = tabulate(resultTable, headers=['Folds', "1","2","3","4","5","6","7","8","9","10","11","12
40 "Classification complete printing results")
41 table)
42 f"Highest mean is {highestMean} so the best K using 5 fold cross validation is {bestAverage
43

```

Saved successfully!

Classification complete printing results

Folds	1	2	3	4	5	6	7	8	9	10
1	0.365	0.37	0.375	0.392	0.394	0.406	0.408	0.416	0.418	0.424
2	0.365	0.37	0.375	0.392	0.394	0.406	0.408	0.416	0.418	0.424
3	0.365	0.37	0.375	0.392	0.394	0.406	0.408	0.416	0.418	0.424
4	0.365	0.37	0.375	0.392	0.394	0.406	0.408	0.416	0.418	0.424
5	0.365	0.37	0.375	0.392	0.394	0.406	0.408	0.416	0.418	0.424
Mean	0.365	0.37	0.375	0.392	0.394	0.406	0.408	0.416	0.418	0.424

Highest mean is 0.442 so the best K using 5 fold cross validation is 18

```

1 #Using best k to classify test database
2 print(f"Classifying using KNN test dataset using best K {bestAverageK}")
3 clf = KNeighborsClassifier(n_neighbors=bestAverageK)
4 clf.fit(X_train_lsa , y_train)
5 knn_accuracy = accuracy_score(y_test, clf.predict(X_test_lsa))
6 print(f"Test dataset accuracy for best K {bestAverageK} is {knn_accuracy} ")
7 print(classification_report(y_test, clf.predict(X_test_lsa), target_names=target_names))
8

```

```

Classifying using KNN test dataset using best K 18
Test dataset accuracy for best K 18 is 0.46436

```

	precision	recall	f1-score	support
Rating 1	0.49	0.73	0.59	7324
Rating 2	0.35	0.22	0.27	5176
Rating 3	0.36	0.27	0.31	5157
Rating 4	0.54	0.51	0.52	7343
accuracy			0.46	25000
macro avg	0.43	0.43	0.42	25000
weighted avg	0.45	0.46	0.45	25000

```

1 print(f"Classifying using multinomialNB")
2 clf = MultinomialNB()
3 clf.fit(X_train_lsa_scaled, y_train)
4 nb_accuracy = accuracy_score(y_test, clf.predict(X_test_lsa_scaled))
5 print(f"Test dataset accuracy is {nb_accuracy} ")
6 print(classification_report(y_test, clf.predict(X_test_lsa_scaled), target_names=target_names))

```

Saved successfully!

```

Classifying using multinomialNB
Test dataset accuracy is 0.4126

```

	precision	recall	f1-score	support
Rating 1	0.36	0.98	0.52	7324
Rating 2	0.00	0.00	0.00	5176
Rating 3	0.00	0.00	0.00	5157
Rating 4	0.64	0.43	0.51	7343
accuracy			0.41	25000
macro avg	0.25	0.35	0.26	25000
weighted avg	0.29	0.41	0.30	25000

```

/usr/local/lib/python3.6/dist-packages/sklearn/metrics/_classification.py:1272: UndefinedWarning:
    _warn_prf(average, modifier, msg_start, len(result))

```

```

1 print(f"Classifying using SVM")
2 clf = svm.SVC()
3 clf.fit(X_train_lsa_scaled , y_train)
4 svm_accuracy = accuracy_score(y_test, clf.predict(X_test_lsa_scaled))
5 print(f"Test dataset accuracy is {svm_accuracy} ")
6 print(classification_report(y_test, clf.predict(X_test_lsa_scaled), target_names=target_names))

```

```
6 print(classification_report(y_test, clf.predict(X_test_15d_scaled), target_names=target_names))
7
```

Classifying using SVM

Test dataset accuracy is 0.56256

	precision	recall	f1-score	support
Rating 1	0.59	0.79	0.68	7324
Rating 2	0.45	0.32	0.37	5176
Rating 3	0.50	0.24	0.33	5157
Rating 4	0.59	0.73	0.66	7343
accuracy			0.56	25000
macro avg	0.53	0.52	0.51	25000
weighted avg	0.54	0.56	0.54	25000

Saved successfully!

